

**12.10 Suppose you need to sort a relation of 40 gigabytes, with 4 kilobyte blocks, using a memory size of 40 megabytes. Suppose the cost of a seek is 5 milliseconds, while the disk transfer rate is 40 megabytes per second.**

**a. Find the cost of sorting the relation, in seconds, with  $bb = 1$  and with  $bb = 100$ .**

**Ans:**  $br = 40 \text{ GB} / 4 \text{ KB} = 10000000 \text{ blocks}$

$M = 40 \text{ MB} / 4 \text{ KB} = 10000 \text{ blocks}$

The initial number of runs =  $(br / M) = 1000$

**b. In each case, how many merge passes are required?**

**Ans:** The number of merge passes required =  $(\log_{M-1} (br / M)) = (\log_{9999} 1000) = 1$

Block transfers =  $br (2 * 1 + 1) = 30000000 \text{ block}$

Seeks =  $2 (br / M) + (br / bb) (2 * 1 - 1)$

Therefore, if  $bb = 12000 + 10000000 = 10002000 \text{ seeks}$

$bb = 1002000 + 100000 = 102000 \text{ seeks}$

Total sorting cost in seconds =  $(\text{hash}(\#) \text{ of block transfers}) * 4\text{KB} / 40\text{MB} + (\text{hash}(\#) \text{ of seeks}) * 5/1000$

**c. Suppose a flash storage device is used instead of a disk, and it has a seek time of 1 microsecond, and a transfer rate of 40 megabytes per second. Recompute the cost of sorting the relation, in seconds, with  $bb = 1$  and with  $bb = 100$ , in this setting.**

**Ans:** if  $bb = 130000000 * 4 \text{ KB} / 40 \text{ MB} + 10002000 * 5 / 1000 = 3000 + 50010 = 53010 \text{ sec}$

$bb = 10030000000 * 4 \text{ KB} / 40 \text{ MB} + 102000 * 5 / 1000 = 3000 + 510 = 3510 \text{ sec}$

**12.12 Why is it not desirable to force users to make an explicit choice of a query processing strategy? Are there cases in which it is desirable for users to be aware of the costs of competing query-processing strategies? Explain your answer.**

**Ans:** In general it is not desirable to force users to choose a query processing strategy because naive users might select an inefficient strategy. The reason users would make poor choices about processing queries is that they would not know how a relation is stored, nor about its indices. It is unreasonable to force users to be aware of these details since ease of use is a major object of database query languages. If users are aware of the costs of different strategies, they could write queries efficiently, thus helping performance. This could happen if experts were using the system.

**12.13 Design a variant of the hybrid merge-join algorithm for the case where both relations are not physically sorted, but both have a sorted secondary index on the join attributes.**

**Ans:** We merge the leaf entries of the first sorted secondary index with the leaf entries of the second sorted secondary index. The result file contains pairs of addresses, the first address in each pair pointing to a tuple in the first relation, and the second address pointing to a tuple in the second relation.

This result file is first sorted on the first relation's addresses. The relation is then scanned in physical storage order and addresses in the result file are replaced by the actual tuple values. Then the result file is sorted on the second relation's addresses, allowing a scan of the second relation in physical storage order to complete the join

**12.14 Estimate the number of blocks transfers and seeks required by your solution to Exercise 12.13 for  $r_1 \bowtie r_2$ , where  $r_1$  and  $r_2$  are as defined in Practice Exercise 12.3.**

**Ans:**  $r_1$  occupies 800 blocks, and  $r_2$  occupies 1500 blocks. Let there be  $n$  pointers per index leaf block (we assume that both the indices have leaf blocks and pointers of equal sizes). Let us assume  $M$  pages of memory,  $M < 800$ .  $r_1$ 's index will need  $B_1 = (20000/n)$  leaf blocks, and  $r_2$ 's index will need  $B_2 = (45000/n)$  leaf blocks. Therefore, the merge join will need  $B_3 = B_1 + B_2$  accesses, without output. The number of output tuples is estimated as  $no = (20000 * 45000 / \max(V(C, r_1), V(C, r_2)))$ . Each output tuple will need two pointers, so the number of blocks of join output will be  $Bo_1 = (no * n / 2)$ . Hence the join needs  $B_j = B_3 + Bo_1$  disk block accesses. Now we must replace the pointers by actual tuples. For the first sorting,  $B_{s1} = Bo_1(2(\log M - 1(B_{o1}/M)) + 2)$  disk accesses are needed, including the writing of output to disk. The number of blocks of  $r_1$  which must be accessed in order to replace the pointers with tuple values is  $\min(800, no)$ . Let  $n_1$  pairs of the form ( $r_1$  tuple, pointer to  $r_2$ ) fit in one disk block. Therefore, the intermediate result after replacing the  $r_1$  pointers will occupy  $Bo_2 = (no/n_1)$  Blocks. Hence the first pass of replacing the  $r_1$ -pointers will cost  $B_f = B_{s1} + Bo_1 + \min(800, no) + Bo_2$  disk accesses. The second pass for replacing the  $r_2$ -pointers has a similar analysis. Let  $n_2$  tuples of the final join fit in one block. Then the second pass of replacing the  $r_2$ -pointers will cost  $B_s = B_{s2} + Bo_2 + \min(1500, no)$  disk accesses, where  $B_{s2} = Bo_2(2(\log M - 1(B_{o2}/M)) + 2)$ . Hence the total number of disk accesses for the join is  $B_j + B_f + B_s$ , and the number of pages of output is  $(no/n_2)$ .

**12.15 The hash-join algorithm as described in Section 12.5.5 computes the natural join of two relations. Describe how to extend the hash-join algorithm to compute the natural left outer join, the natural right outer join and the natural full outer join. (Hint: Keep extra information with each tuple in the hash index, to detect whether any tuple in the probe relation matches the tuple in the hash index.) Try out your algorithm on the takes and student relations.**

**Ans:** For the probe relation tuple  $tr$  under consideration, if no matching tuple is found in the build relation's hash partition, it is padded with nulls and included in the result. This will give us the natural left outer join  $tr \bowtie_{\text{left}} ts$ . To get the natural right outer join  $tr \bowtie_{\text{right}} ts$ , we can keep a Boolean flag with each tuple in the current build relation partition  $H_{si}$  residing in memory and set it whenever any probe relation tuple matches with it. When we are finished with  $H_{si}$ , all the tuples in it which do not have their flag set, are padded with nulls and included in the result. To get the natural full outer join, we do both the above operations together. To try out our algorithm, we use the sample customer and depositor relations of Figures 13.17 and 13.18. Let us assume that there is enough memory to hold three tuples of the build relation plus a hash index for those three tuples. We use depositor as the build relation. We use the simple hashing function which returns the first letter of customer-name. Taking the first partitions, we get  $H_{r1} = \{("Adams", "Spring", "Pittsfield")\}$ , and  $H_{s1} = \phi$ . The tuple in the probe relation partition will have no matching tuple, so  $(("Adams", "Spring", "Pittsfield", \text{null}))$  is outputted. In the partition for "D", the lone build relation tuple is unmatched, thus giving an output tuple  $(("David", \text{null}, \text{null}, A-306))$ . In the partition for "H", we find a match for the first time, producing the output tuple  $(("Hayes", "Main", "Harrison", A-102))$ . Proceeding in a similar way, we process all the partitions and complete the join.

**12.18 Suppose you have to compute  $AGsum(C)(r)$  as well as  $ABGsum(C)(r)$ . Describe how to compute these together using a single sorting of  $r$ .**

**Ans:** The estimated size of the relation can be determined by calculating the average number of tuples which would be joined with each tuple of the second relation. In this case, for each tuple in  $r_1$ ,  $1500/V(C, r_2) = 15/11$  tuples (on the average) of  $r_2$  would join with it. The intermediate relation would have  $15000/11$  tuples. This relation is joined with  $r_3$  to yield a result of approximately 10,227 tuples ( $15000/11 \times 750/100 = 10227$ ). A good strategy should join  $r_1$  and  $r_2$  first, since the intermediate relation is about the same size as  $r_1$  or  $r_2$ . Then  $r_3$  is joined to this result.